

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

02.Load the Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.



Read the Dataset

```
mydata=pd.read_csv('/content/drive/MyDrive/Colab Notebooks/abalone.csv')
```

```
mydata.shape
```

(4177, 9)

```
mydata.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
mydata.tail()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

```
mydata.columns
```

```
Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
       'Viscera weight', 'Shell weight', 'Rings'],
      dtype='object')
```

```
mydata.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	

```
mydata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings            4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```


```
mydata.dtypes
```

```
Sex              object
Length           float64
Diameter         float64
Height           float64
Whole weight     float64
Shucked weight   float64
Viscera weight   float64
Shell weight     float64
Rings            int64
dtype: object
```

### 03.Perform Virtualization


```
sns.boxplot(mydata[ 'Rings' ])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c649e0d0>
```




```
sns.boxplot(mydata[ 'Length' ])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c647c410>
```





```
sns.boxplot(mydata[ 'Diameter' ])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c63f23d0>
```





```
sns.boxplot(mydata['Whole weight'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c63d15d0>
```




```
sns.boxplot(mydata['Shucked weight'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c6336250>
```




```
sns.boxplot(mydata['Viscera weight'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c62b7510>
```




```
sns.boxplot(mydata['Shell weight'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas  
FutureWarning  
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c6290bd0>
```




```
sns.boxplot(mydata['Rings'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c61f2910>
```



```
sns.heatmap(mydata.corr(),annot=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f30c6184490>
```

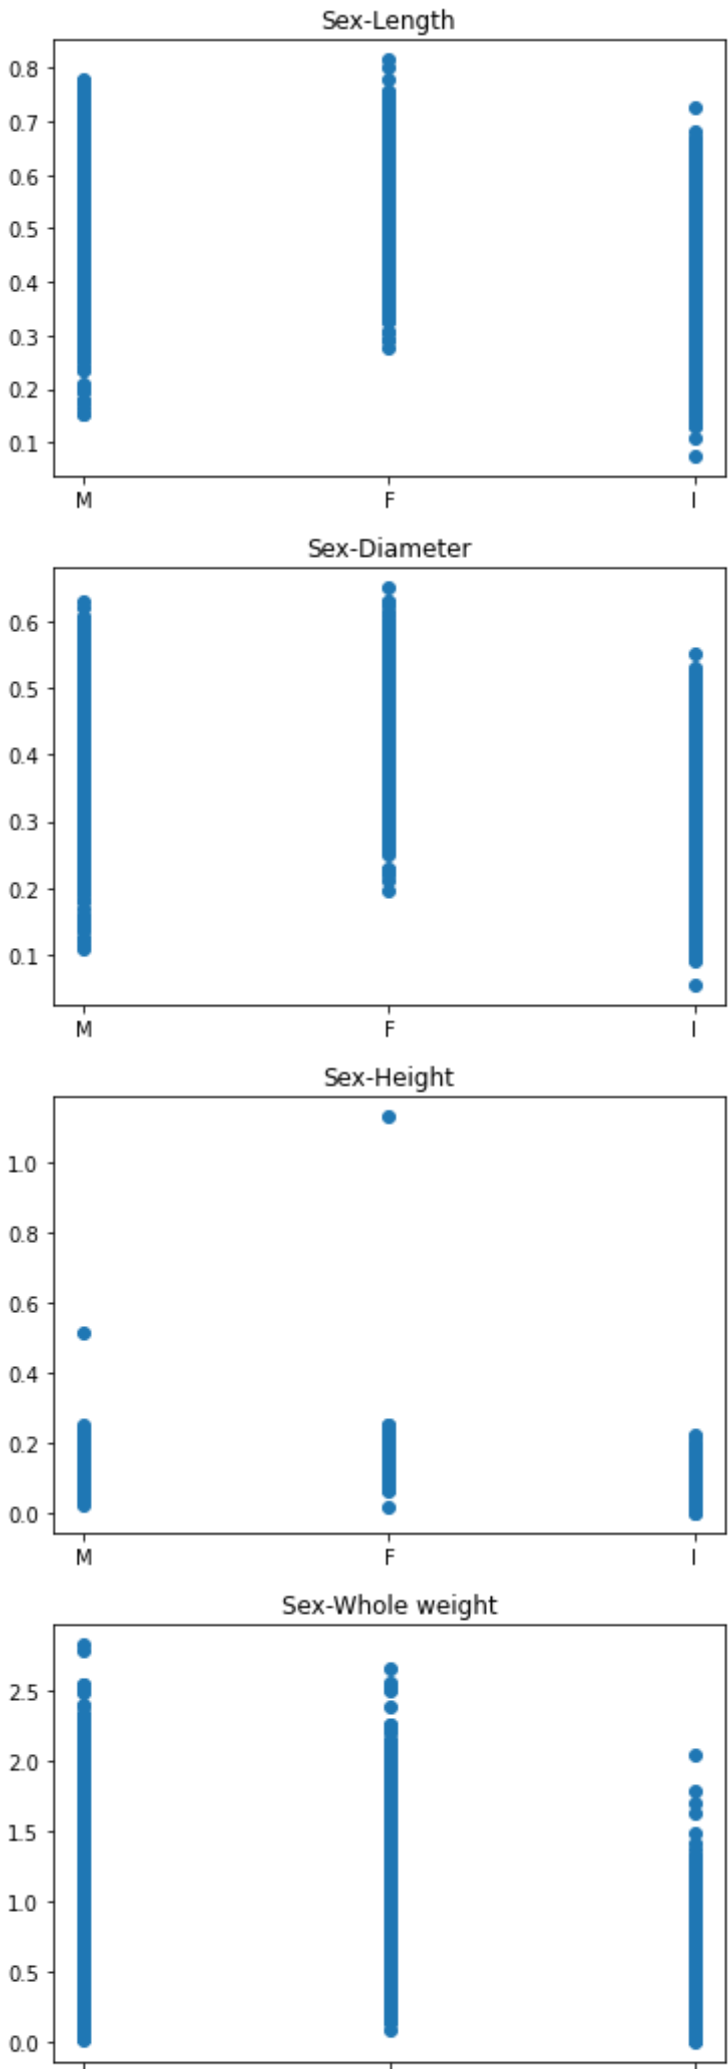


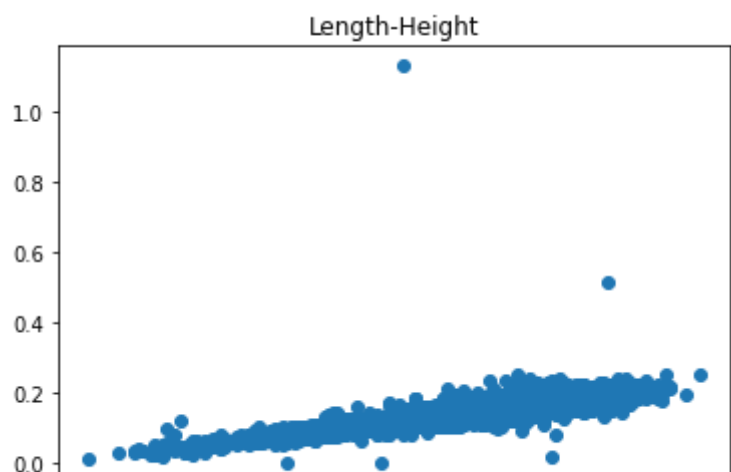


```
mydata.iloc[:, -4:-1].sum().sum(), len(mydata)
```

```
(3253.014, 4177)
```

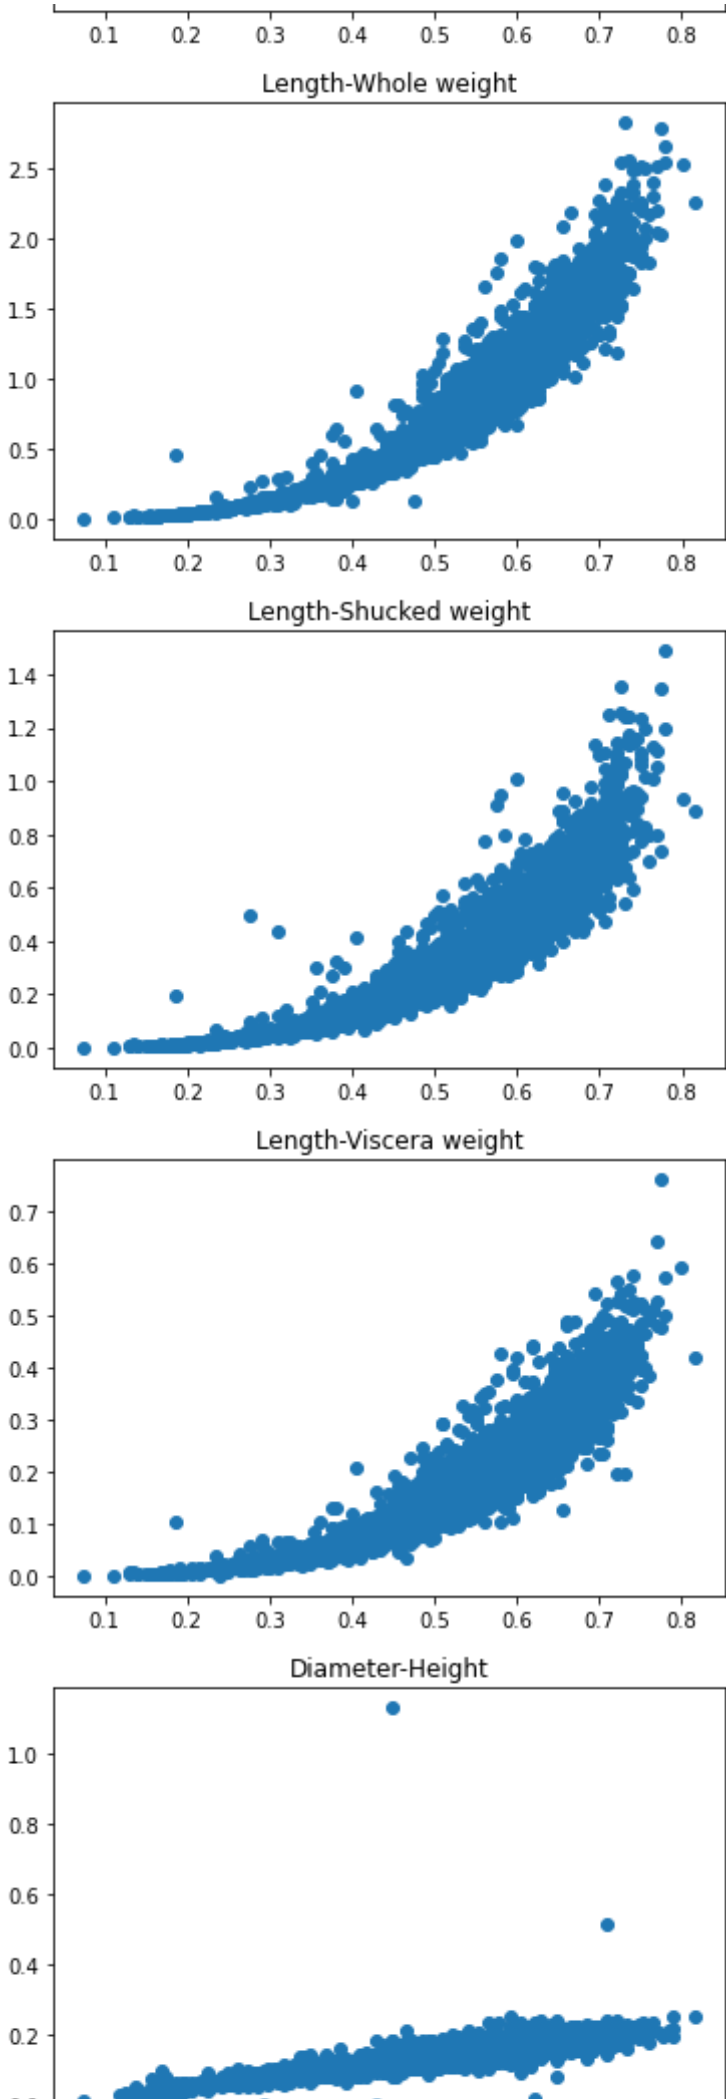
## Bivariate Analysis


```
cols = list(mydata.iloc[:, :7].columns)
for i in range(len(cols)-1):
    for j in range(i+1, len(cols)):
        plt.scatter(mydata[cols[i]], mydata[cols[j]])
        plt.title(cols[i]+'-'+cols[j])
        plt.show()
```














Diameter-Whole weight




Diameter-Shucked weight

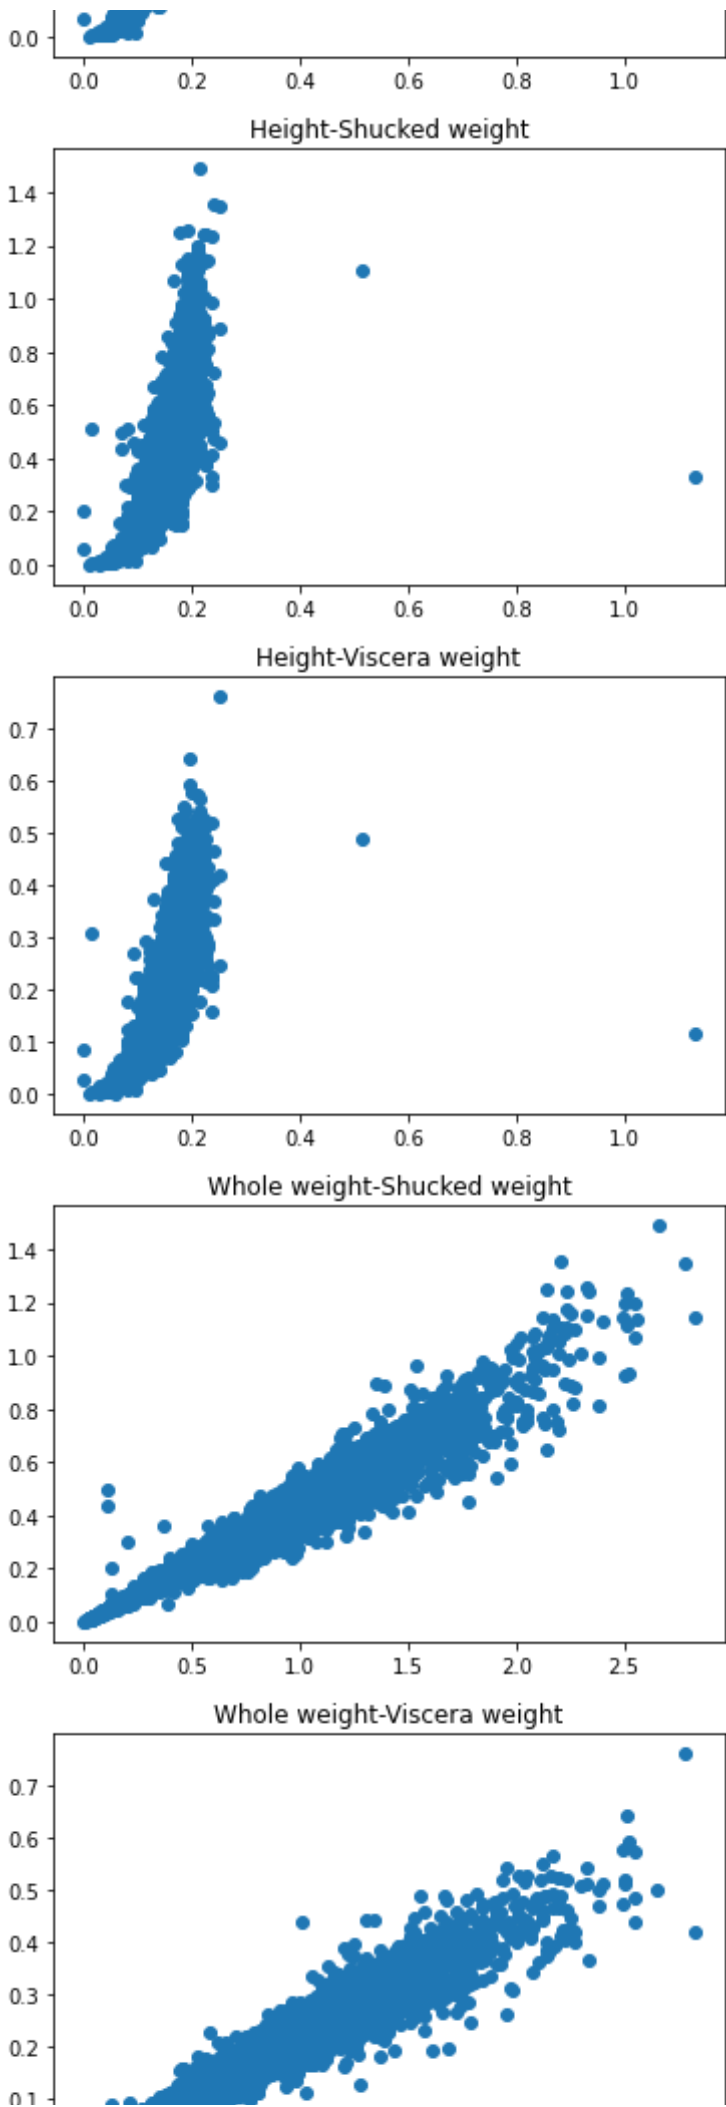



Diameter-Viscera weight



Height-Whole weight







#### 04.Perform Descriptive Statistics

```
mydata = pd.DataFrame()
mydata.sum()
```

Sex	MMFMIIFFMFFMFMFFMIFMMMIFFFFFMMMMFMFFMFFMFFMFFIIII...
Length	2188.715
Diameter	1703.72
Height	582.76
Whole weight	3461.656
Shucked weight	1501.078
Viscera weight	754.3395
Shell weight	997.5965
Rings	41493
dtype: object	

```
mydata.sum(1)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropp
    """Entry point for launching an IPython kernel.
0      16.9045
1       8.1485
2      11.3700
3      11.9305
4       8.0540
...
4172    13.9250
4173    13.0450
4174    12.5770
4175    13.4425
4176    17.2255
Length: 4177, dtype: float64
```

```
mydata.mean()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropp
    """Entry point for launching an IPython kernel.
Length      0.523992
Diameter     0.407881
Height       0.139516
Whole weight  0.828742
Shucked weight 0.359367
Viscera weight 0.180594
Shell weight  0.238831
```

```
Rings          9.933684
dtype: float64
```

```
mydata.std()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropp
    """Entry point for launching an IPython kernel.
Length          0.120093
Diameter        0.099240
Height          0.041827
Whole weight    0.490389
Shucked weight  0.221963
Viscera weight  0.109614
Shell weight    0.139203
Rings           3.224169
dtype: float64
```

```
mydata.count()
```

```
Sex            4177
Length         4177
Diameter       4177
Height         4177
Whole weight   4177
Shucked weight 4177
Viscera weight 4177
Shell weight   4177
Rings          4177
dtype: int64
```

```
mydata.min()
```

```
Sex            F
Length         0.075
Diameter       0.055
Height         0.0
Whole weight   0.002
Shucked weight 0.001
Viscera weight 0.0005
Shell weight   0.0015
Rings          1
dtype: object
```

```
mydata.describe
```

```
<bound method NDFrame.describe of
Shucked weight \
0      M    0.455    0.365    0.095    0.5140    0.2245
1      M    0.350    0.265    0.090    0.2255    0.0995
2      F    0.530    0.420    0.135    0.6770    0.2565
3      M    0.440    0.365    0.125    0.5160    0.2155
4      I    0.330    0.255    0.080    0.2050    0.0895
...  ..    ...    ...    ...    ...    ...
4172  F    0.565    0.450    0.165    0.8870    0.3700
4173  M    0.590    0.440    0.135    0.9660    0.4390
```

4174	M	0.600	0.475	0.205	1.1760	0.5255
4175	F	0.625	0.485	0.150	1.0945	0.5310
4176	M	0.710	0.555	0.195	1.9485	0.9455

	Viscera weight	Shell weight	Rings
0	0.1010	0.1500	15
1	0.0485	0.0700	7
2	0.1415	0.2100	9
3	0.1140	0.1550	10
4	0.0395	0.0550	7
...	...	...	...
4172	0.2390	0.2490	11
4173	0.2145	0.2605	10
4174	0.2875	0.3080	9
4175	0.2610	0.2960	10
4176	0.3765	0.4950	12

```
[4177 rows x 9 columns]>
```

## 05.Handling Missig Values

```
mydata.duplicated().sum()
```

```
0
```

```
mydata.isna().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight
Viscera weight
Shell weight 0
Rings        0
dtype: int64
```

```
mydata.nunique()
```

```
Sex          3
Length      134
Diameter    111
Height      51
Whole weight 2429
Shucked weight
Viscera weight
Shell weight 926
Rings       28
dtype: int64
```

```
mydata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings            4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB

```

```
mydata.drop(columns=['Whole weight','Shucked weight','Viscera weight','Shell weight']).des
```

	Length	Diameter	Height	Rings
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000
<b>mean</b>	0.523992	0.407881	0.139516	9.933684
<b>std</b>	0.120093	0.099240	0.041827	3.224169
<b>min</b>	0.075000	0.055000	0.000000	1.000000
<b>25%</b>	0.450000	0.350000	0.115000	8.000000
<b>50%</b>	0.545000	0.425000	0.140000	9.000000
<b>75%</b>	0.615000	0.480000	0.165000	11.000000
<b>max</b>	0.815000	0.650000	1.130000	29.000000

```
qnt=mydata.drop(columns=['Shucked weight','Viscera weight','Shell weight'])
```

## 06.Find Outliers

```

qnt=mydata.drop(columns=['Sex','Viscera weight','Shucked weight']).quantile(q=[0.015,0.050,0.950])
qnt

```

	Length	Diameter	Height	Whole weight	Shell weight	Rings
<b>0.015</b>	0.215	0.1582	0.050	0.04932	0.0150	4.0
<b>0.050</b>	0.295	0.2200	0.075	0.12590	0.0384	6.0
<b>0.080</b>	0.335	0.2500	0.080	0.17954	0.0550	6.0

```

Q1=qnt.iloc[0]
Q4=qnt.iloc[1]
Q7=qnt.iloc[2]

```

```
iqr=Q4-Q1  
iqr
```

```
Length      0.08000  
Diameter    0.06180  
Height      0.02500  
Whole weight 0.07658  
Shell weight 0.02340  
Rings       2.00000  
dtype: float64
```

```
iqr=Q7-Q1  
iqr
```

```
Length      0.12000  
Diameter    0.09180  
Height      0.03000  
Whole weight 0.13022  
Shell weight 0.04000  
Rings       2.00000  
dtype: float64
```

```
upper=qnt.iloc[2]+1.5*iqr  
upper
```

```
Length      0.51500  
Diameter    0.38770  
Height      0.12500  
Whole weight 0.37487  
Shell weight 0.11500  
Rings       9.00000  
dtype: float64
```

```
lower=qnt.iloc[0]-1.5*iqr  
lower
```

```
Length      0.03500  
Diameter    0.02050  
Height      0.00500  
Whole weight -0.14601  
Shell weight -0.04500  
Rings       1.00000  
dtype: float64
```

```
medium=qnt.iloc[1]-1.5*iqr  
medium
```

```
Length      0.11500  
Diameter    0.08230  
Height      0.03000  
Whole weight -0.06943  
Shell weight -0.02160  
Rings       3.00000  
dtype: float64
```



## Replace Outliers

```
mydata['Rings'] = np.where(mydata['Rings'] > 11.2, 9.933684, mydata['Rings'])
mydata['Whole weight'] = np.where(mydata['Whole weight'] > 3.1, 2.825500, mydata['Whole weight'])
```

## 07.Categorical Columns

```
mydata['Sex'].replace({'M': 1, 'F': 0, 'I': 2}, inplace=True)
mydata.head(10)
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	9.933684
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.000000
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.000000
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.000000
4	2	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.000000
5	2	0.425	0.300	0.095	0.3515	0.1410	0.0775	0.120	8.000000
6	0	0.530	0.415	0.150	0.7775	0.2370	0.1415	0.330	9.933684
7	0	0.545	0.425	0.125	0.7680	0.2940	0.1495	0.260	9.933684
8	1	0.475	0.370	0.125	0.5095	0.2165	0.1125	0.165	9.000000
9	0	0.550	0.440	0.150	0.8945	0.3145	0.1510	0.320	9.933684

## Perform Encoding

```
mydata_all = mydata.drop(columns="Rings")
target = mydata['Rings']
```

```
mydata_all = pd.get_dummies(mydata_all)
mydata_all.columns
```

```
Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight'],
      dtype='object')
```

```
mydata_all.shape
```

```
(4177, 8)
```

```
target.value_counts()
```

```
9.933684    960
9.000000    689
```

```

10.000000    634
8.000000     568
11.000000    487
7.000000     391
6.000000     259
5.000000     115
4.000000      57
3.000000      15
1.000000       1
2.000000       1
Name: Rings, dtype: int64

```

## Dropping Unwanted Columns

```

mydata =mydata.drop(columns=['Length'])
mydata.head()

```

	Sex	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	1	0.365	0.095	0.5140	0.2245	0.1010	0.150	9.933684
1	1	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.000000
2	0	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.000000
3	1	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.000000
4	2	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.000000

## 08.Split the Data into Depenent and Indepenent Variable

```

Y= mydata['Rings']
mydata = mydata.drop(['Rings'], axis = 1)
Y= mydata

```

```

X=mydata.iloc[:, :-1]
X.head()

```

	Sex	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	1	0.365	0.095	0.5140	0.2245	0.1010
1	1	0.265	0.090	0.2255	0.0995	0.0485
2	0	0.420	0.135	0.6770	0.2565	0.1415
3	1	0.365	0.125	0.5160	0.2155	0.1140
4	2	0.255	0.080	0.2050	0.0895	0.0395

```

Y=mydata.iloc[:, -1]
Y.head()

```

```

0    0.150

```

```
1    0.070
2    0.210
3    0.155
4    0.055
```

```
Name: Shell weight, dtype: float64
```

## 09.Scale The Independent Variables

```
from sklearn.preprocessing import StandardScaler
```

```
cls=StandardScaler()
X=cls.fit_transform(X)
```

```
X
```

```
array([[ -0.0105225 , -0.43214879, -1.06442415, -0.64189823, -0.60768536,
        -0.72621157],
       [ -0.0105225 , -1.439929   , -1.18397831, -1.23027711, -1.17090984,
        -1.20522124],
       [ -1.26630752,  0.12213032, -0.10799087, -0.30946926, -0.4634999 ,
        -0.35668983],
       ...,
       [ -0.0105225 ,  0.67640943,  1.56576738,  0.70821206,  0.74855917,
         0.97541324],
       [ -1.26630752,  0.77718745,  0.25067161,  0.54199757,  0.77334105,
         0.73362741],
       [ -0.0105225 ,  1.48263359,  1.32665906,  2.28368063,  2.64099341,
         1.78744868]])
```

## 10.Split Data Into Training and Testing

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,Y_train,Y_test= train_test_split(X,Y,test_size=0.4,random_state=0)
print('train data points :', len(X_train))
print('test data points :', len(X_test))
```

```
train data points : 2506
test data points : 1671
```

```
X_train.shape
```

```
(2506, 6)
```

```
X_test.shape
```

```
(1671, 6)
```

```
X_train
```

```
array([[ 1.24526253, -2.14537514, -1.78174911, -1.47399037, -1.37592355,
        -1.46525506],
       [-0.0105225 , -0.73448285, -0.34709919, -0.77752109, -0.64373173,
        -0.49811173],
       [-0.0105225 ,  0.82757646,  0.72888826,  0.7051529 ,  0.74855917,
        0.84311534],
       ...,
       [-0.0105225 ,  0.42446438,  0.13111745,  0.26565325,  0.46694694,
        0.23636976],
       [-1.26630752,  0.82757646,  0.6093341 ,  0.60827942,  0.53002808,
        0.51008957],
       [ 1.24526253, -0.83526087, -0.70576167, -1.02531323, -1.02221858,
        -0.96343541]])
```

X\_test

```
array([[ -0.0105225 ,  0.21659075,  0.17251933, ...,  0.18101643,
        -0.36887819,  0.56939553],
       [ 1.24526253, -0.1998034 , -0.07942572, ..., -0.43387519,
        -0.44322382, -0.34300384],
       [-0.0105225 ,  0.79954256,  0.72679844, ...,  0.87034766,
        0.75531787,  1.7646387 ],
       ...,
       [ 1.24526253,  0.67462432,  0.62602042, ...,  0.22486442,
        -0.09402464,  0.18618779],
       [-0.0105225 ,  0.46642724,  0.47485339, ..., -0.06779544,
        0.20561078, -0.12402799],
       [ 1.24526253, -1.61554351, -1.69187405, ..., -1.3577422 ,
        -1.28355474, -1.34664314]])
```

## 11. Build the Model

### 1. Linear Regression

```
from sklearn.linear_model import LinearRegression
```

```
lm = LinearRegression()
lm.fit(X_train, Y_train)
```

```
LinearRegression()
```

```
Y_train_pred = lm.predict(X_train)
Y_test_pred = lm.predict(X_test)
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
s = mean_squared_error(Y_train, Y_train_pred)
print('Mean Squared error of training set :%2f'%s)
```

```
p = mean_squared_error(Y_test, Y_test_pred)
print('Mean Squared error of testing set :%2f'%p)
```

```
Mean Squared error of training set :0.000995
Mean Squared error of testing set :0.000791
```

```
from sklearn.metrics import r2_score
s = r2_score(Y_train, Y_train_pred)
print('R2 Score of training set: %.2f'%s)
```

```
p = r2_score(Y_test, Y_test_pred)
print('R2 Score of testing set: %.2f'%p)
```

```
R2 Score of training set:0.95
R2 Score of testing set:0.96
```

## 2.Ridge

```
from sklearn.linear_model import Ridge
```

```
ridge_mod = Ridge(alpha=0.01, normalize=True)
ridge_mod.fit(X_train, Y_train)
ridge_mod.fit(X_test, Y_test)
ridge_model_pred = ridge_mod.predict(X_test)
ridge_mod.score(X_train, Y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_base.py:145: FutureWarning
If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), Ridge())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```

Set parameter alpha to: `original_alpha * n_samples`.

```
FutureWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_base.py:145: FutureWarning
If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), Ridge())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```


Set parameter alpha to: `original_alpha * n_samples`.

```
FutureWarning,
0.9290023156453548
```

```
ridge_mod.score(X_test, Y_test)
```

```
0.9519921522913208
```

```
plt.scatter(Y_test, ridge_model_pred)
plt.xlabel('True Values')
plt.ylabel('Predictions')
```



### 3.Support Vector Regression

```
from sklearn.svm import SVR
```

```
svr = SVR(kernel = 'linear')
svr.fit(X_train, Y_train)
svr.fit(X_test, Y_test)
```

```
SVR(kernel='linear')
```

```
Y_train_pred = svr.predict(X_train)
Y_test_pred = svr.predict(X_test)
svr.score(X_train, Y_train)
```

```
0.8922092465754603
```

```
svr.score(X_test, Y_test)
```

```
0.93330364267072
```

### 4.Random Forest Regression

```
from sklearn.ensemble import RandomForestRegressor
```